


# CS/DS 541: Class 13

Jacob Whitehill

# Recurrent neural networks (RNNs)

# Recurrent neural network


- We can construct a simple **recurrent neural network** (RNN) as follows:



$$\begin{aligned}\hat{y}_t &= g(\mathbf{x}_1, \dots, \mathbf{x}_t; \mathbf{U}, \mathbf{V}, \mathbf{w}) &= \mathbf{h}_t^\top \mathbf{w} \\ \mathbf{h}_t &= \sigma(\mathbf{U}\mathbf{h}_{t-1} + \mathbf{V}\mathbf{x}_t)\end{aligned}$$


# Difficulty in training RNNs

- In their simplest form, RNNs are typically hard to train:
- The gradients can occasionally become very large (**exploding gradient**), which forces us to use a very small learning rate (which makes training slow).




# Difficulty in training RNNs

- In their simplest form, RNNs are typically hard to train:
- The gradients can also become very small (**vanishing gradient**), which also makes learning very slow.



# Difficulty in training RNNs

- A related problem is that, if  $T$  is large, then information *early* in the input sequence (e.g.,  $\mathbf{x}_1$ ) can “get lost” when trying to predict values *late* in the sequence (e.g.,  $\hat{y}_T$ ).



# Difficulty in training deep FFNNs

- Another strategy for preventing vanishing and exploding gradients is to use **skip connections** (more later).
  - These are used in LSTM and GRU RNNs, as well as ResNet FFNNs.
- Yet another strategy is to restrict **U** to the manifold of **unitary matrices** (i.e., all eigenvalues have magnitude 1; see Helfrich & Ye 2019).

# Long short-term memory (LSTM) neural networks



# LSTMs

- <https://colah.github.io/posts/2015-08-Understanding-LSTMs/>

# LSTMs

- Three gates — forget (f), input (i), and output (o).

$$\mathbf{f}_t = \sigma(\mathbf{W}_f[\mathbf{x}_t, \mathbf{h}_{t-1}] + \mathbf{b}_f)$$

$$\mathbf{i}_t = \sigma(\mathbf{W}_i[\mathbf{x}_t, \mathbf{h}_{t-1}] + \mathbf{b}_i)$$

$$\mathbf{o}_t = \sigma(\mathbf{W}_o[\mathbf{x}_t, \mathbf{h}_{t-1}] + \mathbf{b}_o)$$

# LSTMs

- Three gates — forget (f), input (i), and output (o).
- Two state vectors:  $\mathbf{h}_t$ ,  $\mathbf{c}_t$ .

$$\mathbf{f}_t = \sigma(\mathbf{W}_f[\mathbf{x}_t, \mathbf{h}_{t-1}] + \mathbf{b}_f)$$

$$\mathbf{i}_t = \sigma(\mathbf{W}_i[\mathbf{x}_t, \mathbf{h}_{t-1}] + \mathbf{b}_i)$$

$$\mathbf{o}_t = \sigma(\mathbf{W}_o[\mathbf{x}_t, \mathbf{h}_{t-1}] + \mathbf{b}_o)$$

$$\tilde{\mathbf{c}}_t = \tanh(\mathbf{W}_c[\mathbf{x}_t, \mathbf{h}_{t-1}] + \mathbf{b}_c)$$

$$\mathbf{c}_t = \mathbf{f}_t \odot \mathbf{c}_{t-1} + \mathbf{i}_t \odot \tilde{\mathbf{c}}_t$$

$$\mathbf{h}_t = \mathbf{o}_t \odot \tanh(\mathbf{c}_t)$$

# LSTMs

- In total, we have 4 weight matrices and 4 bias vectors.

$$\mathbf{f}_t = \sigma(\mathbf{W}_f[\mathbf{x}_t, \mathbf{h}_{t-1}] + \mathbf{b}_f)$$

$$\mathbf{i}_t = \sigma(\mathbf{W}_i[\mathbf{x}_t, \mathbf{h}_{t-1}] + \mathbf{b}_i)$$

$$\mathbf{o}_t = \sigma(\mathbf{W}_o[\mathbf{x}_t, \mathbf{h}_{t-1}] + \mathbf{b}_o)$$

$$\tilde{\mathbf{c}}_t = \tanh(\mathbf{W}_c[\mathbf{x}_t, \mathbf{h}_{t-1}] + \mathbf{b}_c)$$

$$\mathbf{c}_t = \mathbf{f}_t \odot \mathbf{c}_{t-1} + \mathbf{i}_t \odot \tilde{\mathbf{c}}_t$$


$$\mathbf{h}_t = \mathbf{o}_t \odot \tanh(\mathbf{c}_t)$$

# LSTM

- The memory cell  $\mathbf{c}$  offers a pathway through the network to preserve information across long time-spans:

$$\mathbf{c}_t = \mathbf{i}_t \odot \tilde{\mathbf{c}}_t + \mathbf{f}_t \odot \mathbf{c}_{t-1}$$


- It tends not to decay due to exponentiated eigenvalues.



# LSTM

- If  $\mathbf{f}_t=1$ , then  $\mathbf{c}_t$  directly contains information from 1, ...,  $t$ :


$$\mathbf{c}_t = \mathbf{i}_t \odot \tilde{\mathbf{c}}_t + \mathbf{c}_{t-1}$$



# LSTM

- If  $\mathbf{f}_t=1$ , then  $\mathbf{c}_t$  directly contains information from 1, ...,  $t$ :

$$\begin{aligned}\mathbf{c}_t &= \mathbf{i}_t \odot \tilde{\mathbf{c}}_t + \mathbf{c}_{t-1} \\ &= \mathbf{i}_t \odot \tilde{\mathbf{c}}_t + \mathbf{i}_{t-1} \odot \tilde{\mathbf{c}}_{t-1} + \mathbf{c}_{t-2}\end{aligned}$$




# LSTM

- If  $\mathbf{f}_t=1$ , then  $\mathbf{c}_t$  directly contains information from 1, ...,  $t$ :

$$\mathbf{c}_t = \mathbf{i}_t \odot \tilde{\mathbf{c}}_t + \mathbf{c}_{t-1}$$

$$= \mathbf{i}_t \odot \tilde{\mathbf{c}}_t + \mathbf{i}_{t-1} \odot \tilde{\mathbf{c}}_{t-1} + \mathbf{c}_{t-2}$$

$$= \mathbf{i}_t \odot \tilde{\mathbf{c}}_t + \mathbf{i}_{t-1} \odot \tilde{\mathbf{c}}_{t-1} + \mathbf{i}_{t-2} \odot \tilde{\mathbf{c}}_{t-2} + \mathbf{c}_{t-3}$$





# LSTM

- If  $\mathbf{f}_t=1$ , then  $\mathbf{c}_t$  directly contains information from 1, ...,  $t$ :


$$\mathbf{c}_t = \mathbf{i}_t \odot \tilde{\mathbf{c}}_t + \mathbf{c}_{t-1}$$

$$= \mathbf{i}_t \odot \tilde{\mathbf{c}}_t + \mathbf{i}_{t-1} \odot \tilde{\mathbf{c}}_{t-1} + \mathbf{c}_{t-2}$$

$$= \mathbf{i}_t \odot \tilde{\mathbf{c}}_t + \mathbf{i}_{t-1} \odot \tilde{\mathbf{c}}_{t-1} + \mathbf{i}_{t-2} \odot \tilde{\mathbf{c}}_{t-2} + \mathbf{c}_{t-3}$$

...

$$\mathbf{c}_t = \mathbf{i}_t \odot \tilde{\mathbf{c}}_t + \mathbf{i}_{t-1} \odot \tilde{\mathbf{c}}_{t-1} + \dots + \mathbf{i}_2 \odot \tilde{\mathbf{c}}_2 + \mathbf{c}_1$$



# LSTM

- If  $\mathbf{f}_t=1$ , then  $\mathbf{c}_t$  directly contains information from 1, ...,  $t$ :


$$\mathbf{c}_t = \mathbf{i}_t \odot \tilde{\mathbf{c}}_t + \mathbf{c}_{t-1}$$

$$= \mathbf{i}_t \odot \tilde{\mathbf{c}}_t + \mathbf{i}_{t-1} \odot \tilde{\mathbf{c}}_{t-1} + \mathbf{c}_{t-2}$$

$$= \mathbf{i}_t \odot \tilde{\mathbf{c}}_t + \mathbf{i}_{t-1} \odot \tilde{\mathbf{c}}_{t-1} + \mathbf{i}_{t-2} \odot \tilde{\mathbf{c}}_{t-2} + \mathbf{c}_{t-3}$$


...

$$\mathbf{c}_t = \mathbf{i}_t \odot \tilde{\mathbf{c}}_t + \mathbf{i}_{t-1} \odot \tilde{\mathbf{c}}_{t-1} + \dots + \mathbf{i}_2 \odot \tilde{\mathbf{c}}_2 + \mathbf{c}_1$$



# LSTM

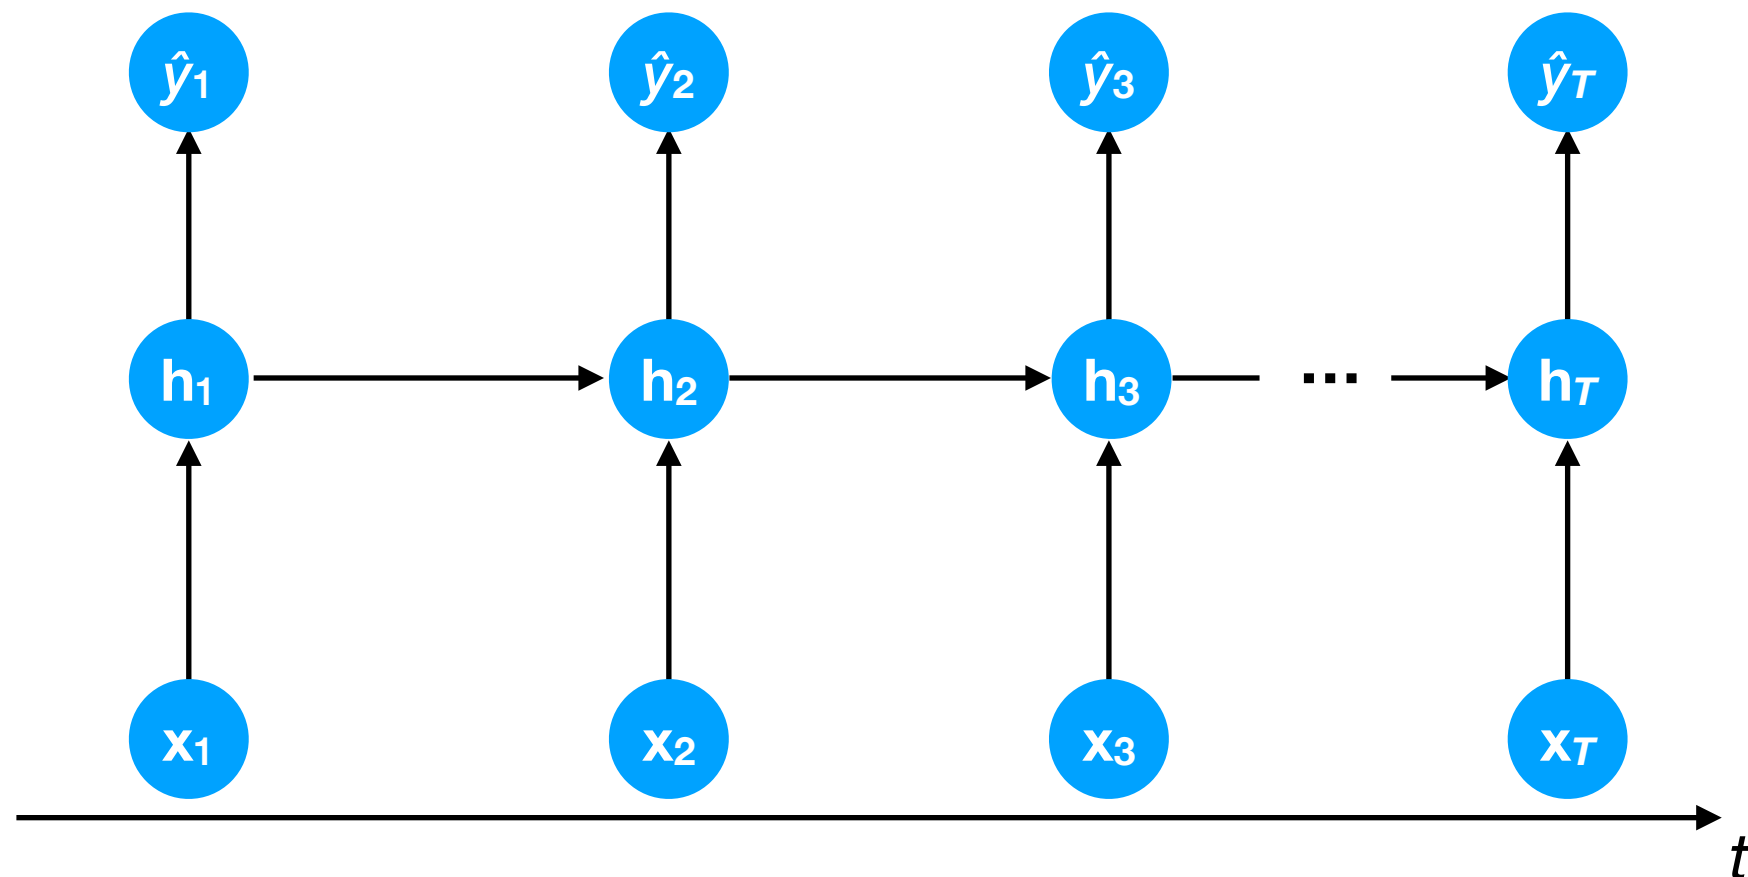
- This is sometimes called a **skip connection**.



# Bi-directional RNNs


# Bi-directional RNNs

- The RNNs (including LSTMs) we have examined so far are useful when an output  $\hat{y}_t$  must be estimated immediately after  $t$  timesteps.




# Bi-directional RNNs

- However, in some settings, we may be able to wait to see the *entire* input sequence before producing any outputs.
- In this case, it can help to harness the entire sequence  $x_1, \dots, x_T$  for each prediction  $\hat{y}_t$ .



# Bi-directional RNNs

- With a bi-directional RNN, each prediction  $\hat{y}_t$  is determined by two different hidden representations  $\mathbf{h}_t, \mathbf{h}'_t$  — one from each direction of processing.



# Regularization



# Regularization

- We (like in Goodfellow's *Deep Learning*) can define **regularization** as anything that helps to improve generalization performance of a trained ML model.
- Deep learning benefits from many standard techniques (e.g.,  $L_1$ ,  $L_2$  regularization) but also offers some of its own.

# $L_2$ regularization

- We have already seen  $L_2$  regularization, whereby the  $L_2$  norm of each (vectorized) weight matrix is added to the loss, e.g.:

$$f(\hat{\mathbf{y}}, \mathbf{y}; \{\mathbf{W}^{(k)}, \mathbf{b}^{(k)}\}_{k=1}^l) = \frac{1}{2}(\hat{\mathbf{y}} - \mathbf{y})^\top (\hat{\mathbf{y}} - \mathbf{y}) + \sum_{k=1}^l \frac{\alpha_k}{2} \|\mathbf{W}^{(k)}\|_{\text{Fr}}^2$$

- The Frobenius norm of a matrix is the sum of its **squared** entries; it is equivalent to the  $L_2$  norm of the vectorized matrix. Gradient:  $\nabla_{\mathbf{W}} \left( \frac{1}{2} \|\mathbf{W}\|_{\text{Fr}}^2 \right) = \mathbf{W}$
- The  $L_2$  norm encourages *all* the entries in each weight matrix to be small.

# $L_2$ regularization

- We can apply different amounts of regularization to each matrix  $\mathbf{W}^{(k)}$ .

$$f(\hat{\mathbf{y}}, \mathbf{y}; \{\mathbf{W}^{(k)}, \mathbf{b}^{(k)}\}_{k=1}^l) = \frac{1}{2}(\hat{\mathbf{y}} - \mathbf{y})^\top (\hat{\mathbf{y}} - \mathbf{y}) + \sum_{k=1}^l \frac{\alpha_k}{2} \|\mathbf{W}^{(k)}\|_{\text{Fr}}^2$$

- Bias terms are typically not regularized because we want them to “shift” the activations as much as needed.

# $L_1$ regularization

- A related technique is  $L_1$  regularization, which penalizes the sum of the **absolute values** of each entry of a weight matrix, e.g.:

$$f(\hat{\mathbf{y}}, \mathbf{y}; \{\mathbf{W}^{(k)}, \mathbf{b}^{(k)}\}_{k=1}^l) = \frac{1}{2}(\hat{\mathbf{y}} - \mathbf{y})^\top (\hat{\mathbf{y}} - \mathbf{y}) + \sum_{k=1}^l \alpha_k \|\mathbf{W}^{(k)}\|_1$$

- The  $L_1$  norm encourages *some* parameters to be *exactly* 0. This can encourage sparse feature representations.

Gradient:

$$\nabla_{\mathbf{W}} (\|\mathbf{W}\|_1) = \text{sign}(\mathbf{W})$$

- Note that  $L_1$  and  $L_2$  regularization can also be combined.

# $L_2$ regularization = weight decay

- You may sometimes encounter the term **weight decay**, which means that weights tend to “decay” in magnitude during training.
- Weight decay is equivalent to  $L_2$  regularization in SGD:

$$\mathbf{W}^{\text{new}} = \mathbf{W} - \epsilon(\nabla_{\mathbf{W}} f(\mathbf{W}) + \alpha \mathbf{W})$$

# $L_2$ regularization = weight decay

- You may sometimes encounter the term **weight decay**, which means that weights tend to “decay” in magnitude during training.
- Weight decay is equivalent to  $L_2$  regularization in SGD:

$$\begin{aligned}\mathbf{W}^{\text{new}} &= \mathbf{W} - \epsilon(\nabla_{\mathbf{W}} f(\mathbf{W}) + \alpha \mathbf{W}) \\ &= \mathbf{W} - \epsilon \nabla_{\mathbf{W}} f(\mathbf{W}) - \epsilon \alpha \mathbf{W}\end{aligned}$$

# $L_2$ regularization = weight decay

- You may sometimes encounter the term **weight decay**, which means that weights tend to “decay” in magnitude during training.
- Weight decay is equivalent to  $L_2$  regularization in SGD:

$$\begin{aligned}\mathbf{W}^{\text{new}} &= \mathbf{W} - \epsilon(\nabla_{\mathbf{W}} f(\mathbf{W}) + \alpha \mathbf{W}) \\ &= \mathbf{W} - \epsilon \nabla_{\mathbf{W}} f(\mathbf{W}) - \epsilon \alpha \mathbf{W} \\ &= \mathbf{W}(1 - \epsilon \alpha) - \epsilon \nabla_{\mathbf{W}} f(\mathbf{W})\end{aligned}$$

# $L_2$ regularization = weight decay

- You may sometimes encounter the term **weight decay**, which means that weights tend to “decay” in magnitude during training.
- Weight decay is equivalent to  $L_2$  regularization in SGD:

$$\begin{aligned}\mathbf{W}^{\text{new}} &= \mathbf{W} - \epsilon(\nabla_{\mathbf{W}} f(\mathbf{W}) + \alpha \mathbf{W}) \\ &= \mathbf{W} - \epsilon \nabla_{\mathbf{W}} f(\mathbf{W}) - \epsilon \alpha \mathbf{W} \\ &= \mathbf{W}(1 - \epsilon \alpha) - \epsilon \nabla_{\mathbf{W}} f(\mathbf{W})\end{aligned}$$

- For  $\epsilon \alpha < 1$ ,  $\mathbf{W}$  shrinks in length at each iteration.



# $L_2$ regularization $\cong$ Gaussian noise augmentation

- For 2-layer linear NNs (i.e., linear regression),  $L_2$  regularization is also equivalent to augmenting the training set by adding element-wise Gaussian noise to each input.
- To show this, we will use a probabilistic interpretation.
- Let  $\mathbf{x} \in \mathbb{R}^m$  be a randomly drawn training input and (scalar)  $y$  is its associated label.
- Let  $\mathbf{n} \in \mathbb{R}^m$ ,  $\mathbf{n} \sim \mathcal{N}(\mathbf{0}, \alpha \mathbf{I})$  be 0-mean Gaussian noise that is *independent* of  $\mathbf{x}$ .
- Recall that, for any two independent random variables  $\mathbf{x}$  and  $\mathbf{n}$ , we have:  $\mathbb{E}[\mathbf{x}\mathbf{n}] = \mathbb{E}[\mathbf{x}]\mathbb{E}[\mathbf{n}]$
- Define  $\tilde{\mathbf{x}} = \mathbf{x} + \mathbf{n}$ .

# $L_2$ regularization $\cong$ Gaussian noise augmentation

- Instead of a sum, the cost function contains the *expected*  $L_2$  distance between the predictions and target labels.
- We define separate cost functions for the original ( $\mathbf{x}$ ) and the noise-augmented ( $\tilde{\mathbf{x}}$ ) inputs w.r.t. weights  $\mathbf{w}$ :

$$f(\mathbf{w}) = \mathbb{E}[(\hat{y} - y)^2] = \mathbb{E}[(\mathbf{x}^\top \mathbf{w} - y)^2]$$

$$\tilde{f}(\mathbf{w}) = \mathbb{E}[(\tilde{\hat{y}} - y)^2] = \mathbb{E}[(\tilde{\mathbf{x}}^\top \mathbf{w} - y)^2]$$

# $L_2$ regularization $\cong$ Gaussian noise augmentation

- We can then derive that the noise-augmented loss  $\tilde{f}$  equals the original loss  $f$ , plus an  $L_2$  regularization term:

$$\mathbb{E}[(\tilde{\mathbf{x}}^\top \mathbf{w} - y)^2] = \mathbb{E}[(\mathbf{x} + \mathbf{n})^\top \mathbf{w} - y]^2]$$

# $L_2$ regularization $\cong$ Gaussian noise augmentation

- We can then derive that the noise-augmented loss  $\tilde{f}$  equals the original loss  $f$ , plus an  $L_2$  regularization term:

$$\begin{aligned}\mathbb{E}[(\tilde{\mathbf{x}}^\top \mathbf{w} - y)^2] &= \mathbb{E}[(\mathbf{x} + \mathbf{n})^\top \mathbf{w} - y]^2 \\ &= \mathbb{E}[(\mathbf{x}^\top \mathbf{w} + \mathbf{n}^\top \mathbf{w} - y)^2]\end{aligned}$$

# $L_2$ regularization $\cong$ Gaussian noise augmentation

- We can then derive that the noise-augmented loss  $\tilde{f}$  equals the original loss  $f$ , plus an  $L_2$  regularization term:

$$\begin{aligned}\mathbb{E}[(\tilde{\mathbf{x}}^\top \mathbf{w} - y)^2] &= \mathbb{E}[((\mathbf{x} + \mathbf{n})^\top \mathbf{w} - y)^2] \\ &= \mathbb{E}[(\mathbf{x}^\top \mathbf{w} + \mathbf{n}^\top \mathbf{w} - y)^2] \\ &= \mathbb{E}[((\mathbf{x}^\top \mathbf{w} - y) + \mathbf{n}^\top \mathbf{w})^2]\end{aligned}$$

# $L_2$ regularization $\cong$ Gaussian noise augmentation

- We can then derive that the noise-augmented loss  $\tilde{f}$  equals the original loss  $f$ , plus an  $L_2$  regularization term:

$$\begin{aligned}\mathbb{E}[(\tilde{\mathbf{x}}^\top \mathbf{w} - y)^2] &= \mathbb{E}[((\mathbf{x} + \mathbf{n})^\top \mathbf{w} - y)^2] \\ &= \mathbb{E}[(\mathbf{x}^\top \mathbf{w} + \mathbf{n}^\top \mathbf{w} - y)^2] \\ &= \mathbb{E}[((\mathbf{x}^\top \mathbf{w} - y) + \mathbf{n}^\top \mathbf{w})^2] \\ &= \mathbb{E}[(\mathbf{x}^\top \mathbf{w} - y)^2] + \mathbb{E}[2(\mathbf{x}^\top \mathbf{w} - y)\mathbf{n}^\top \mathbf{w}] + \mathbb{E}[(\mathbf{n}^\top \mathbf{w})^2]\end{aligned}$$

# $L_2$ regularization $\cong$ Gaussian noise augmentation

- We can then derive that the noise-augmented loss  $\tilde{f}$  equals the original loss  $f$ , plus an  $L_2$  regularization term:

$$\begin{aligned}\mathbb{E}[(\tilde{\mathbf{x}}^\top \mathbf{w} - y)^2] &= \mathbb{E}[(\mathbf{x} + \mathbf{n})^\top \mathbf{w} - y]^2 \\ &= \mathbb{E}[(\mathbf{x}^\top \mathbf{w} + \mathbf{n}^\top \mathbf{w} - y)^2] \\ &= \mathbb{E}[(\mathbf{x}^\top \mathbf{w} - y) + \mathbf{n}^\top \mathbf{w}]^2 \\ &= \mathbb{E}[(\mathbf{x}^\top \mathbf{w} - y)^2] + \mathbb{E}[2(\mathbf{x}^\top \mathbf{w} - y)\mathbf{n}^\top \mathbf{w}] + \mathbb{E}[(\mathbf{n}^\top \mathbf{w})^2] \\ &= f(\mathbf{w}) + \mathbb{E}[2(\mathbf{x}^\top \mathbf{w} \mathbf{n}^\top \mathbf{w}) - 2y\mathbf{n}^\top \mathbf{w}] + \mathbb{E}[\mathbf{w}^\top \mathbf{n} \mathbf{n}^\top \mathbf{w}]\end{aligned}$$

# $L_2$ regularization $\cong$ Gaussian noise augmentation

- We can then derive that the noise-augmented loss  $\tilde{f}$  equals the original loss  $f$ , plus an  $L_2$  regularization term:

$$\begin{aligned}\mathbb{E}[(\tilde{\mathbf{x}}^\top \mathbf{w} - y)^2] &= \mathbb{E}[(\mathbf{x} + \mathbf{n})^\top \mathbf{w} - y]^2 \\&= \mathbb{E}[(\mathbf{x}^\top \mathbf{w} + \mathbf{n}^\top \mathbf{w} - y)^2] \\&= \mathbb{E}[(\mathbf{x}^\top \mathbf{w} - y) + \mathbf{n}^\top \mathbf{w}]^2 \\&= \mathbb{E}[(\mathbf{x}^\top \mathbf{w} - y)^2] + \mathbb{E}[2(\mathbf{x}^\top \mathbf{w} - y)\mathbf{n}^\top \mathbf{w}] + \mathbb{E}[(\mathbf{n}^\top \mathbf{w})^2] \\&= f(\mathbf{w}) + \mathbb{E}[2(\mathbf{x}^\top \mathbf{w} \mathbf{n}^\top \mathbf{w}) - 2y \mathbf{n}^\top \mathbf{w}] + \mathbb{E}[\mathbf{w}^\top \mathbf{n} \mathbf{n}^\top \mathbf{w}] \\&= f(\mathbf{w}) + 2\mathbb{E}[\mathbf{x}^\top] \mathbf{w} \mathbb{E}[\mathbf{n}^\top] \mathbf{w} - 2\mathbb{E}[y] \mathbb{E}[\mathbf{n}^\top] \mathbf{w} + \mathbf{w}^\top \mathbb{E}[\mathbf{n} \mathbf{n}^\top] \mathbf{w}\end{aligned}$$

Here we can split the expectation into the product of two expectations.



# $L_2$ regularization $\cong$ Gaussian noise augmentation

- We can then derive that the noise-augmented loss  $\tilde{f}$  equals the original loss  $f$ , plus an  $L_2$  regularization term:

$$\begin{aligned}\mathbb{E}[(\tilde{\mathbf{x}}^\top \mathbf{w} - y)^2] &= \mathbb{E}[((\mathbf{x} + \mathbf{n})^\top \mathbf{w} - y)^2] \\&= \mathbb{E}[(\mathbf{x}^\top \mathbf{w} + \mathbf{n}^\top \mathbf{w} - y)^2] \\&= \mathbb{E}[((\mathbf{x}^\top \mathbf{w} - y) + \mathbf{n}^\top \mathbf{w})^2] \\&= \mathbb{E}[(\mathbf{x}^\top \mathbf{w} - y)^2] + \mathbb{E}[2(\mathbf{x}^\top \mathbf{w} - y)\mathbf{n}^\top \mathbf{w}] + \mathbb{E}[(\mathbf{n}^\top \mathbf{w})^2] \\&= f(\mathbf{w}) + \mathbb{E}[2(\mathbf{x}^\top \mathbf{w}\mathbf{n}^\top \mathbf{w}) - 2y\mathbf{n}^\top \mathbf{w}] + \mathbb{E}[\mathbf{w}^\top \mathbf{n}\mathbf{n}^\top \mathbf{w}] \\&= f(\mathbf{w}) + 2\mathbb{E}[\mathbf{x}^\top] \mathbf{w} \mathbb{E}[\mathbf{n}^\top] \mathbf{w} - 2\mathbb{E}[y] \mathbb{E}[\mathbf{n}^\top] \mathbf{w} + \mathbf{w}^\top \mathbb{E}[\mathbf{n}\mathbf{n}^\top] \mathbf{w} \\&= f(\mathbf{w}) + 2\mathbb{E}[\mathbf{x}^\top] \mathbf{w} \cdot \mathbf{0} \cdot \mathbf{w} - 2\mathbb{E}[y] \cdot \mathbf{0} \cdot \mathbf{w} + \mathbf{w}^\top \alpha \mathbf{I} \mathbf{w}\end{aligned}$$

$\mathbf{n}$  has 0 mean.

# $L_2$ regularization $\cong$ Gaussian noise augmentation

- We can then derive that the noise-augmented loss  $\tilde{f}$  equals the original loss  $f$ , plus an  $L_2$  regularization term:

$$\begin{aligned}\mathbb{E}[(\tilde{\mathbf{x}}^\top \mathbf{w} - y)^2] &= \mathbb{E}[((\mathbf{x} + \mathbf{n})^\top \mathbf{w} - y)^2] \\&= \mathbb{E}[(\mathbf{x}^\top \mathbf{w} + \mathbf{n}^\top \mathbf{w} - y)^2] \\&= \mathbb{E}[((\mathbf{x}^\top \mathbf{w} - y) + \mathbf{n}^\top \mathbf{w})^2] \\&= \mathbb{E}[(\mathbf{x}^\top \mathbf{w} - y)^2] + \mathbb{E}[2(\mathbf{x}^\top \mathbf{w} - y)\mathbf{n}^\top \mathbf{w}] + \mathbb{E}[(\mathbf{n}^\top \mathbf{w})^2] \\&= f(\mathbf{w}) + \mathbb{E}[2(\mathbf{x}^\top \mathbf{w} \mathbf{n}^\top \mathbf{w}) - 2y\mathbf{n}^\top \mathbf{w}] + \mathbb{E}[\mathbf{w}^\top \mathbf{n} \mathbf{n}^\top \mathbf{w}] \\&= f(\mathbf{w}) + 2\mathbb{E}[\mathbf{x}^\top] \mathbf{w} \mathbb{E}[\mathbf{n}^\top] \mathbf{w} - 2\mathbb{E}[y] \mathbb{E}[\mathbf{n}^\top] \mathbf{w} + \mathbf{w}^\top \mathbb{E}[\mathbf{n} \mathbf{n}^\top] \mathbf{w} \\&= f(\mathbf{w}) + 2\mathbb{E}[\mathbf{x}^\top] \mathbf{w} \cdot 0 \cdot \mathbf{w} - 2\mathbb{E}[y] \cdot 0 \cdot \mathbf{w} + \mathbf{w}^\top \alpha \mathbf{I} \mathbf{w}\end{aligned}$$

Covariance of  $\mathbf{n}$ .

# $L_2$ regularization $\cong$ Gaussian noise augmentation

- We can then derive that the noise-augmented loss  $\tilde{f}$  equals the original loss  $f$ , plus an  $L_2$  regularization term:

$$\begin{aligned}\mathbb{E}[(\tilde{\mathbf{x}}^\top \mathbf{w} - y)^2] &= \mathbb{E}[((\mathbf{x} + \mathbf{n})^\top \mathbf{w} - y)^2] \\&= \mathbb{E}[(\mathbf{x}^\top \mathbf{w} + \mathbf{n}^\top \mathbf{w} - y)^2] \\&= \mathbb{E}[((\mathbf{x}^\top \mathbf{w} - y) + \mathbf{n}^\top \mathbf{w})^2] \\&= \mathbb{E}[(\mathbf{x}^\top \mathbf{w} - y)^2] + \mathbb{E}[2(\mathbf{x}^\top \mathbf{w} - y)\mathbf{n}^\top \mathbf{w}] + \mathbb{E}[(\mathbf{n}^\top \mathbf{w})^2] \\&= f(\mathbf{w}) + \mathbb{E}[2(\mathbf{x}^\top \mathbf{w}\mathbf{n}^\top \mathbf{w}) - 2y\mathbf{n}^\top \mathbf{w}] + \mathbb{E}[\mathbf{w}^\top \mathbf{n}\mathbf{n}^\top \mathbf{w}] \\&= f(\mathbf{w}) + 2\mathbb{E}[\mathbf{x}^\top] \mathbf{w} \mathbb{E}[\mathbf{n}^\top] \mathbf{w} - 2\mathbb{E}[y] \mathbb{E}[\mathbf{n}^\top] \mathbf{w} + \mathbf{w}^\top \mathbb{E}[\mathbf{n}\mathbf{n}^\top] \mathbf{w} \\&= f(\mathbf{w}) + 2\mathbb{E}[\mathbf{x}^\top] \mathbf{w} \cdot 0 \cdot \mathbf{w} - 2\mathbb{E}[y] \cdot 0 \cdot \mathbf{w} + \mathbf{w}^\top \alpha \mathbf{I} \mathbf{w} \\&= f(\mathbf{w}) + \alpha \mathbf{w}^\top \mathbf{w}\end{aligned}$$

$L_2$  regularization 43



# $L_2$ regularization $\cong$ Gaussian noise augmentation

- For non-linear and deep NNs, element-wise Gaussian noise augmentation and  $L_2$  regularization are no longer equivalent, but they may sometimes have similar effects.

# Weight sharing

- One of the most powerful methods of regularizing a neural network is to reduce the number of parameters by tying some weight matrices to be the same.
- Prominent cases:
  - CNNs: the same convolution filter is used at *every location*.
  - RNNs (including LSTMs, GRUs, etc): the same weights are used at *every timestep*.

# Pre-training

- Both supervised and unsupervised pre-training allow ML practitioners to harness much larger datasets to learn good representations of the inputs.
- For domains with a small number of training data, this can be a powerful regularization technique to prevent overfitting.

# Ensembles

- By training multiple predictors and averaging their outputs, we can create an **ensemble**.
- Ensembles are an easy and often effective way of increasing accuracy.
- Useful ensembles require the individual predictors' outputs to have low correlation, i.e., they make different kinds of mistakes on the same inputs.

# Ensembles

- Suppose we train an ensemble of  $n$  NNs, each of which is unbiased, i.e.,  $E[\hat{y}_i - y] = E[\varepsilon_i] = 0$ , where:
  - $y$  is the target label for a randomly drawn example  $\mathbf{x}$ .
  - $\hat{y}_i$  is the  $i^{\text{th}}$  NN's prediction for  $\mathbf{x}$ .
  - $\varepsilon_i$  is the error of NN  $i$ 's prediction.
- Let the variance (expected squared error) of NN  $i$  be  $E[\varepsilon_i^2] = v$ .
- Suppose the covariance of predictions between NNs  $i \neq j$  is  $E[\varepsilon_i \varepsilon_j] = c$ .
- The ensemble's prediction on any  $\mathbf{x}$  is:  $\frac{1}{n} \sum_{i=1}^n \hat{y}_i$



# Ensembles

- We can compute the expected squared error of the ensemble:

$$\mathbb{E} \left[ \left( \left( \frac{1}{n} \sum_{i=1}^n \hat{y}_i \right) - y \right)^2 \right] = \mathbb{E} \left[ \left( \frac{1}{n} \sum_{i=1}^n (\hat{y}_i - y) \right)^2 \right]$$

# Ensembles

- We can compute the expected squared error of the ensemble:

$$\begin{aligned}\mathbb{E} \left[ \left( \left( \frac{1}{n} \sum_{i=1}^n \hat{y}_i \right) - y \right)^2 \right] &= \mathbb{E} \left[ \left( \frac{1}{n} \sum_{i=1}^n (\hat{y}_i - y) \right)^2 \right] \\ &= \mathbb{E} \left[ \left( \frac{1}{n} \sum_{i=1}^n \epsilon_i \right)^2 \right]\end{aligned}$$

# Ensembles

- We can compute the expected squared error of the ensemble:

$$\begin{aligned}\mathbb{E} \left[ \left( \left( \frac{1}{n} \sum_{i=1}^n \hat{y}_i \right) - y \right)^2 \right] &= \mathbb{E} \left[ \left( \frac{1}{n} \sum_{i=1}^n (\hat{y}_i - y) \right)^2 \right] \\ &= \mathbb{E} \left[ \left( \frac{1}{n} \sum_{i=1}^n \epsilon_i \right)^2 \right] \\ &= \frac{1}{n^2} \mathbb{E} \left[ \begin{array}{c} \epsilon_1 \epsilon_1 + \epsilon_1 \epsilon_2 + \dots + \epsilon_1 \epsilon_n + \\ \dots \\ + \epsilon_n \epsilon_1 + \epsilon_n \epsilon_2 + \dots + \epsilon_n \epsilon_n \end{array} \right]\end{aligned}$$

# Ensembles

- We can compute the expected squared error of the ensemble:

$$\begin{aligned}
 \mathbb{E} \left[ \left( \left( \frac{1}{n} \sum_{i=1}^n \hat{y}_i \right) - y \right)^2 \right] &= \mathbb{E} \left[ \left( \frac{1}{n} \sum_{i=1}^n (\hat{y}_i - y) \right)^2 \right] \\
 &= \mathbb{E} \left[ \left( \frac{1}{n} \sum_{i=1}^n \epsilon_i \right)^2 \right] \\
 &= \frac{1}{n^2} \mathbb{E} \left[ \begin{array}{c} \epsilon_1 \epsilon_1 + \epsilon_1 \epsilon_2 + \dots + \epsilon_1 \epsilon_n + \\ \dots \\ + \epsilon_n \epsilon_1 + \epsilon_n \epsilon_2 + \dots + \epsilon_n \epsilon_n \end{array} \right] \\
 &= \frac{1}{n^2} \left[ \begin{array}{c} v + c + \dots + c + \\ \dots \\ + c + \dots + c + v \end{array} \right]
 \end{aligned}$$

# Ensembles

- We can compute the expected squared error of the ensemble:

$$\begin{aligned}
 \mathbb{E} \left[ \left( \left( \frac{1}{n} \sum_{i=1}^n \hat{y}_i \right) - y \right)^2 \right] &= \mathbb{E} \left[ \left( \frac{1}{n} \sum_{i=1}^n (\hat{y}_i - y) \right)^2 \right] \\
 &= \mathbb{E} \left[ \left( \frac{1}{n} \sum_{i=1}^n \epsilon_i \right)^2 \right] \\
 &= \frac{1}{n^2} \mathbb{E} \left[ \begin{array}{c} \epsilon_1 \epsilon_1 + \epsilon_1 \epsilon_2 + \dots + \epsilon_1 \epsilon_n + \\ \dots \\ + \epsilon_n \epsilon_1 + \epsilon_n \epsilon_2 + \dots + \epsilon_n \epsilon_n \end{array} \right] \\
 &= \frac{1}{n^2} \left[ \begin{array}{c} v + c + \dots + c + \\ \dots \\ + c + \dots + c + v \end{array} \right] \\
 &= \frac{nv + (n-1)nc}{n^2}
 \end{aligned}$$

# Ensembles

- We can compute the expected squared error of the ensemble:

$$\begin{aligned}
 \mathbb{E} \left[ \left( \left( \frac{1}{n} \sum_{i=1}^n \hat{y}_i \right) - y \right)^2 \right] &= \mathbb{E} \left[ \left( \frac{1}{n} \sum_{i=1}^n (\hat{y}_i - y) \right)^2 \right] \\
 &= \mathbb{E} \left[ \left( \frac{1}{n} \sum_{i=1}^n \epsilon_i \right)^2 \right] \\
 &= \frac{1}{n^2} \mathbb{E} \left[ \begin{array}{c} \epsilon_1 \epsilon_1 + \epsilon_1 \epsilon_2 + \dots + \epsilon_1 \epsilon_n + \\ \dots \\ + \epsilon_n \epsilon_1 + \epsilon_n \epsilon_2 + \dots + \epsilon_n \epsilon_n \end{array} \right] \\
 &= \frac{1}{n^2} \left[ \begin{array}{c} v + c + \dots + c + \\ \dots \\ + c + \dots + c + v \end{array} \right] \\
 &= \frac{nv + (n-1)nc}{n^2} \\
 &= \frac{v}{n} + \frac{n-1}{n}c
 \end{aligned}$$

# Ensembles

- If  $c=0$  then the ensemble reduces our expected squared error by a factor of  $n$  — great news!

$$\begin{aligned}
 \mathbb{E} \left[ \left( \left( \frac{1}{n} \sum_{i=1}^n \hat{y}_i \right) - y \right)^2 \right] &= \mathbb{E} \left[ \left( \frac{1}{n} \sum_{i=1}^n (\hat{y}_i - y) \right)^2 \right] \\
 &= \mathbb{E} \left[ \left( \frac{1}{n} \sum_{i=1}^n \epsilon_i \right)^2 \right] \\
 &= \frac{1}{n^2} \mathbb{E} \left[ \begin{array}{c} \epsilon_1 \epsilon_1 + \epsilon_1 \epsilon_2 + \dots + \epsilon_1 \epsilon_n + \\ \dots \\ + \epsilon_n \epsilon_1 + \epsilon_n \epsilon_2 + \dots + \epsilon_n \epsilon_n \end{array} \right] \\
 &= \frac{1}{n^2} \left[ \begin{array}{c} v + c + \dots + c + \\ \dots \\ + c + \dots + c + v \end{array} \right] \\
 &= \frac{nv + (n-1)nc}{n^2} \\
 &= \frac{v}{n} + \frac{n-1}{n}c
 \end{aligned}$$

# Ensembles

- If  $c=v$  then the ensemble is no better than any of the individual predictors.

$$\begin{aligned}
 \mathbb{E} \left[ \left( \left( \frac{1}{n} \sum_{i=1}^n \hat{y}_i \right) - y \right)^2 \right] &= \mathbb{E} \left[ \left( \frac{1}{n} \sum_{i=1}^n (\hat{y}_i - y) \right)^2 \right] \\
 &= \mathbb{E} \left[ \left( \frac{1}{n} \sum_{i=1}^n \epsilon_i \right)^2 \right] \\
 &= \frac{1}{n^2} \mathbb{E} \left[ \begin{array}{c} \epsilon_1 \epsilon_1 + \epsilon_1 \epsilon_2 + \dots + \epsilon_1 \epsilon_n + \\ \dots \\ + \epsilon_n \epsilon_1 + \epsilon_n \epsilon_2 + \dots + \epsilon_n \epsilon_n \end{array} \right] \\
 &= \frac{1}{n^2} \left[ \begin{array}{c} v + c + \dots + c + \\ \dots \\ + c + \dots + c + v \end{array} \right] \\
 &= \frac{nv + (n-1)nc}{n^2} \\
 &= \frac{v}{n} + \frac{n-1}{n}c
 \end{aligned}$$




# Dropout

- One of the most recently discovered regularization methods is **dropout**, whereby a random set of neurons is removed from the network at each iteration during both forward and backward propagation.
- Surprisingly, this simple method can both help the network to reach a better local minimum *and* prevent it from overfitting.


# Dropout

- Suppose we are training the NN shown below:




# Dropout

- Suppose we are training the NN shown below:
- For each step of SGD, we randomly select (with “keep” probability  $p$ ) some of the input and hidden neurons (*not* the output neurons).




# Dropout

- Suppose we are training the NN shown below:
- We then remove these neurons and perform forward-propagation on the reduced network.




# Dropout

- Suppose we are training the NN shown below:
- During back-propagation, we adjust the weights of *only* those neurons that were retained in the reduced network.



# Dropout

- We then replace the neurons we had removed and resume training. (During the next SGD iteration, we will randomly select *another* set of neurons to remove, etc.)




# Dropout: example

- Suppose the weights are:

$$\mathbf{W}^{(1)} = \begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \\ 7 & 8 & 9 \end{bmatrix} \quad \mathbf{W}^{(2)} = \begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \end{bmatrix}$$

(For simplicity, assume that  $\mathbf{b}^{(1)}=\mathbf{b}^{(2)}=0$ .)




# Dropout: example

- Suppose the weights are:

$$\mathbf{W}^{(1)} = \begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \\ 7 & 8 & 9 \end{bmatrix} \quad \mathbf{W}^{(2)} = \begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \end{bmatrix}$$

- If we drop the red neurons, then we will obtain  $\hat{\mathbf{y}}=[60, 132]^T$  for the input  $\mathbf{x}=[0, 1, 2]^T$  during forward-propagation.






# Dropout: example

- Suppose the weights are:

$$\mathbf{W}^{(1)} = \begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \\ 7 & 8 & 9 \end{bmatrix} \quad \mathbf{W}^{(2)} = \begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \end{bmatrix}$$

- If we drop the red neurons, then we will obtain  $\hat{\mathbf{y}}=[60, 132]^T$  for the input  $\mathbf{x}=[0, 1, 2]^T$  during forward-propagation.




# Dropout: example

- Suppose the weights are:

$$\mathbf{W}^{(1)} = \begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \\ 7 & 8 & 9 \end{bmatrix} \quad \mathbf{W}^{(2)} = \begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \end{bmatrix}$$

- If we drop the red neurons, then we will obtain  $\hat{\mathbf{y}}=[60, 132]^T$  for the input  $\mathbf{x}=[0, 1, 2]^T$  during forward-propagation.

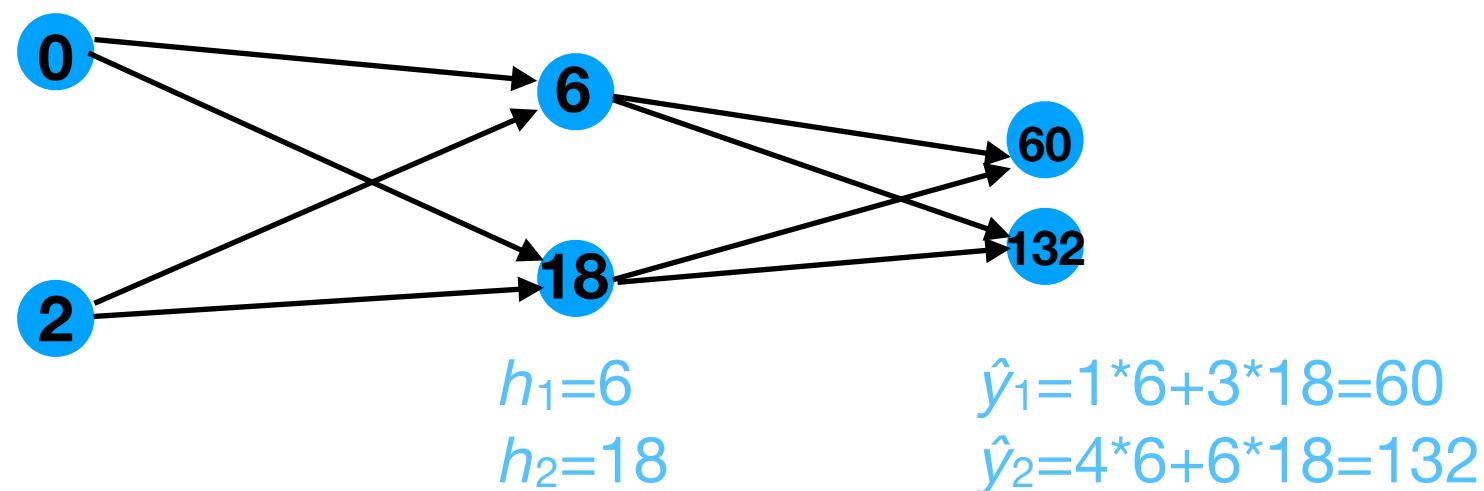


# Dropout: example

- Suppose the weights are:

$$\mathbf{W}^{(1)} = \begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \\ 7 & 8 & 9 \end{bmatrix} \quad \mathbf{W}^{(2)} = \begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \end{bmatrix}$$

- If we drop the red neurons, then we will obtain  $\hat{\mathbf{y}}=[60, 132]^T$  for the input  $\mathbf{x}=[0, 1, 2]^T$  during forward-propagation.




# Dropout: example

- Suppose the weights are:

$$\mathbf{W}^{(1)} = \begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \\ 7 & 8 & 9 \end{bmatrix} \quad \mathbf{W}^{(2)} = \begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \end{bmatrix}$$

- During back-propagation, we will update the weights of *only* those neurons that were not removed.




# Dropout: why helpful?

- There are two main explanations for why dropout helps improve the accuracy of neural networks:
  - Symmetry breaking & prevention of co-adaptation.
  - Ensemble of many smaller networks.

# Symmetry breaking

- When multiple neurons in the hidden layers are highly correlated with each other, the network does not utilize its full capacity.
- Extreme example: all rows of  $\mathbf{W}^{(1)}$  are the same:




- Then all hidden units in  $\mathbf{h}^{(1)}$  are also the same.

# Symmetry breaking

- One reason why we initialize weights randomly is to break symmetry between them, so they learn to produce independent values in the subsequent hidden layer.
- Dropout can also help break symmetry since only some of the elements of each weight matrix are updated during each SGD iteration.

# Co-adaptation


- A subtler problem that can occur is that the weights associated with different “pathways” through the NN can adapt to each other in pathological ways.





# Co-adaptation

- A subtler problem that can occur is that the weights associated with different “pathways” through the NN can adapt to each other in pathological ways.




# Co-adaptation

- If  $\mathbf{W} = \begin{bmatrix} 2 & -3 & 1 & -1 \\ 2 & -3 & -1 & 1 \end{bmatrix}$


then  $\mathbf{h} = \begin{bmatrix} h_1 \\ h_2 \end{bmatrix} = \begin{bmatrix} 2x_1 - 3x_2 + x_3 - x_4 \\ 2x_1 - 3x_2 - x_3 + x_4 \end{bmatrix}$

and hence  $\hat{y} = 2x_1 - 3x_2$ .




# Co-adaptation

- At this point, neither  $x_3$  nor  $x_4$  have any impact on  $\hat{y}$ .
- We may be at a local minimum where the last two columns of  $\mathbf{W}$  stay “locked” to effectively delete  $x_3, x_4$ .
- This is an example of **weight co-adaptation**; it is often a suboptimal solution.




# Co-adaptation

- Instead, we generally want every “pathway” of a NN to give an independently good estimate of  $y$ .
- Dropout can reduce co-adaptation since each pathway must “stand on its own”.



# Ensemble of many smaller networks


- Dropout-based NN training can also be seen as approximating a large ensemble of many smaller networks.
- Each member of the ensemble arises by randomly dropping some of the whole network's neurons:



Ensemble of many networks

# Ensemble of many smaller networks

- At the end of SGD training, the final network approximates the average prediction over all members of the ensemble.
- Caveat: each member of the ensemble is constrained to *share the same weights* with all other members.



Ensemble of many networks