

CS/DS 541: Class 19

Jacob Whitehill

Exercises

https://cs230.stanford.edu/files/cs230exam_fall19_soln.pdf

Exercise 4

Consider a model trying to learn an encoding of some input $x \in \mathbb{R}$. The goal is to encode the input x using $z = w_1 x \in \mathbb{R}$, then accurately reconstruct the original x from the encoded representation using $\hat{x} = w_2 z \in \mathbb{R}$. Here, $(w_1, w_2) \in \mathbb{R} \times \mathbb{R}$. The model is trained with the squared reconstruction error:

$$L(W) = \frac{1}{n} \sum_{i=1}^n (x^{(i)} - w_2 w_1 x^{(i)})^2$$

(2 points) What is the set of solutions for w_1 and w_2 which makes loss zero?

(3 points) Does the loss have a saddle point? Where?

Generative Adversarial Networks (GANs)

Generative Adversarial Networks (GANs)


- However, another entire class of deep learning methods is based on training two networks that **compete against each other** in a zero-sum game.
- In particular, the most prominent method (as of 2020) for generating novel data \mathbf{x} is the Generative Adversarial Network (GAN; Goodfellow et al. 2014).

Generative Adversarial Networks (GANs)

- Like VAEs, GANs consist of two components, but their semantics are different.
- Let $P_{\text{data}}(\mathbf{x})$ be the ground-truth data distribution.
- **Generator G** : given a noise vector \mathbf{z} from an easy-to-sample distribution (e.g., Gaussian, uniform), generate a vector \mathbf{x} that looks like it came from $P_{\text{data}}(\mathbf{x})$.
- **Discriminator D** : given a vector \mathbf{x} , decide if it is real ($\hat{y}=1$) or fake ($\hat{y}=0$). D acts as a “forgery detector”.


Generator G

- Example G with l hidden layers that generates an MNIST image ($28 \times 28 = 784$) \mathbf{x} from a 100-dim noise vector \mathbf{z} :




Discriminator D

- Example D with l hidden layers that estimates $\hat{y} \in (0,1)$ that expresses probability that the input \mathbf{x} is real:




Generative Adversarial Networks (GANs)

- Like VAEs, GANs are trained such that one component “feeds” to the other.
- In contrast to VAEs, the discriminator D is sometimes given a “fake” data vector \mathbf{x} (generated by G), and sometimes given a “real” vector \mathbf{x} sampled from the training set (which approximates $P_{\text{data}}(\mathbf{x})$).



Generative Adversarial Networks (GANs)

- Each network has its own parameters:
 - G has parameters θ_G .
 - D has parameters θ_D .




Generative Adversarial Networks (GANs)

- We can define the following loss on how well D can discriminate fake from real data:

$$f_{\text{acc}}(\theta_G, \theta_D) = \mathbb{E}_{\mathbf{x} \sim P_{\text{data}}(\mathbf{x})} [\log D_{\theta_D}(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim P(\mathbf{z})} [\log(1 - D_{\theta_D}(G_{\theta_G}(\mathbf{z})))]$$

Log-likelihood that D
recognizes real data as real.




Generative Adversarial Networks (GANs)

- We can define the following loss on how well D can discriminate fake from real data:

$$f_{\text{acc}}(\theta_G, \theta_D) = \mathbb{E}_{\mathbf{x} \sim P_{\text{data}}(\mathbf{x})} [\log D_{\theta_D}(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim P(\mathbf{z})} [\log(1 - D_{\theta_D}(G_{\theta_G}(\mathbf{z})))]$$

Log-likelihood that D recognizes fake data as fake.



Generative Adversarial Networks (GANs)

- The goal of D is to *maximize* f_{acc} , whereas the goal of G is to *minimize* f_{acc} .

$$f_{\text{acc}}(\theta_G, \theta_D) = \mathbb{E}_{\mathbf{x} \sim P_{\text{data}}(\mathbf{x})} [\log D_{\theta_D}(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim P(\mathbf{z})} [\log(1 - D_{\theta_D}(G_{\theta_G}(\mathbf{z})))]$$

- This two-player game will reach an equilibrium if we find:

$$\min_{\theta_G} \max_{\theta_D} f_{\text{acc}}(\theta_G, \theta_D)$$

- In particular, this solution corresponds to D having 50% accuracy at detecting forgeries, and G generating fake \mathbf{x} according to $P_{\text{data}}(\mathbf{x})$.

Training GANs

$$f_{\text{acc}}(\theta_G, \theta_D) = \mathbb{E}_{\mathbf{x} \sim P_{\text{data}}(\mathbf{x})} [\log D_{\theta_D}(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim P(\mathbf{z})} [\log(1 - D_{\theta_D}(G_{\theta_G}(\mathbf{z})))]$$

- In practice, we train D and G *iteratively*:
 - Freeze G , and perform SGD on D for k iterations to increase f_{acc} .

Training GANs

$$f_{\text{acc}}(\theta_G, \theta_D) = \mathbb{E}_{\mathbf{x} \sim P_{\text{data}}(\mathbf{x})} [\log D_{\theta_D}(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim P(\mathbf{z})} [\log(1 - D_{\theta_D}(G_{\theta_G}(\mathbf{z})))]$$

- In practice, we train D and G *iteratively*:
- Freeze G , and perform SGD on D for k iterations to increase f_{acc} .

**Improve D 's forgery detection accuracy
for a fixed distribution of fake data.**

Training GANs

$$f_{\text{acc}}(\theta_G, \theta_D) = \mathbb{E}_{\mathbf{x} \sim P_{\text{data}}(\mathbf{x})} [\log D_{\theta_D}(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim P(\mathbf{z})} [\log(1 - D_{\theta_D}(G_{\theta_G}(\mathbf{z})))]$$

- In practice, we train D and G *iteratively*:
 - Freeze G , and perform SGD on D for k iterations to increase f_{acc} .
 - Freeze D , and perform SGD on G for l iterations to decrease f_{acc} .

Training GANs

$$f_{\text{acc}}(\theta_G, \theta_D) = \mathbb{E}_{\mathbf{x} \sim P_{\text{data}}(\mathbf{x})} [\log D_{\theta_D}(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim P(\mathbf{z})} [\log(1 - D_{\theta_D}(G_{\theta_G}(\mathbf{z})))]$$

- In practice, we train D and G *iteratively*:
 - Freeze G , and perform SGD on D for k iterations to increase f_{acc} .
 - Freeze D , and perform SGD on G for l iterations to decrease f_{acc} .

Improve G for a fixed forgery detector D .

Training GANs

Algorithm 1 Minibatch stochastic gradient descent training of generative adversarial nets. The number of steps to apply to the discriminator, k , is a hyperparameter. We used $k = 1$, the least expensive option, in our experiments.

for number of training iterations **do**

for k steps **do**

- Sample minibatch of m noise samples $\{\mathbf{z}^{(1)}, \dots, \mathbf{z}^{(m)}\}$ from noise prior $p_g(\mathbf{z})$.
- Sample minibatch of m examples $\{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(m)}\}$ from data generating distribution $p_{\text{data}}(\mathbf{x})$.
- Update the discriminator by ascending its stochastic gradient:

$$\nabla_{\theta_d} \frac{1}{m} \sum_{i=1}^m \left[\log D(\mathbf{x}^{(i)}) + \log \left(1 - D(G(\mathbf{z}^{(i)})) \right) \right].$$

end for

- Sample minibatch of m noise samples $\{\mathbf{z}^{(1)}, \dots, \mathbf{z}^{(m)}\}$ from noise prior $p_g(\mathbf{z})$.
- Update the generator by descending its stochastic gradient:

$$\nabla_{\theta_g} \frac{1}{m} \sum_{i=1}^m \log \left(1 - D(G(\mathbf{z}^{(i)})) \right).$$

end for

The gradient-based updates can use any standard gradient-based learning rule. We used momentum in our experiments.

Closer look at f_{acc}

- Consider the loss term for fake data:

$$f_{\text{acc}}(\theta_G, \theta_D) = \mathbb{E}_{\mathbf{x} \sim P_{\text{data}}(\mathbf{x})} [\log D_{\theta_D}(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim P(\mathbf{z})} [\log(1 - D_{\theta_D}(G_{\theta_G}(\mathbf{z})))]$$

- What happens early during training, when G is not very good (but D typically is fairly good)?

$$\nabla_{\theta_G} \log(1 - D(G(\mathbf{z}))) = -\frac{1}{1 - D(G(\mathbf{z}))} \frac{\partial D}{\partial G} \frac{\partial G}{\partial \theta_G}$$

Closer look at f_{acc}

- Consider the loss term for fake data:

$$f_{\text{acc}}(\theta_G, \theta_D) = \mathbb{E}_{\mathbf{x} \sim P_{\text{data}}(\mathbf{x})} [\log D_{\theta_D}(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim P(\mathbf{z})} [\log(1 - D_{\theta_D}(G_{\theta_G}(\mathbf{z})))]$$

- What happens early during training, when G is not very good (but D typically is fairly good)?
- D will output ~ 0 .

$$\nabla_{\theta_G} \log(1 - D(G(\mathbf{z}))) = -\frac{1}{1 - D(G(\mathbf{z}))} \frac{\partial D}{\partial G} \frac{\partial G}{\partial \theta_G}$$

Closer look at f_{acc}

- Consider the loss term for fake data:

$$f_{\text{acc}}(\theta_G, \theta_D) = \mathbb{E}_{\mathbf{x} \sim P_{\text{data}}(\mathbf{x})} [\log D_{\theta_D}(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim P(\mathbf{z})} [\log(1 - D_{\theta_D}(G_{\theta_G}(\mathbf{z})))]$$

- What happens early during training, when G is not very good (but D typically is fairly good)?
- D will output ~ 0 .

$$\begin{aligned} \nabla_{\theta_G} \log(1 - D(G(\mathbf{z}))) &= -\frac{1}{1 - D(G(\mathbf{z}))} \frac{\partial D}{\partial G} \frac{\partial G}{\partial \theta_G} \\ &= -\frac{1}{1} \sigma'(v) \frac{\partial v}{\partial G} \frac{\partial G}{\partial \theta_G} \end{aligned}$$

Here we assume D uses a logistic sigmoid σ as its output layer, whose input is v .

Closer look at f_{acc}

- Consider the loss term for fake data:

$$f_{\text{acc}}(\theta_G, \theta_D) = \mathbb{E}_{\mathbf{x} \sim P_{\text{data}}(\mathbf{x})} [\log D_{\theta_D}(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim P(\mathbf{z})} [\log(1 - D_{\theta_D}(G_{\theta_G}(\mathbf{z})))]$$

- What happens early during training, when G is not very good (but D typically is fairly good)?
- D will output ~ 0 .

$$\begin{aligned}\nabla_{\theta_G} \log(1 - D(G(\mathbf{z}))) &= -\frac{1}{1 - D(G(\mathbf{z}))} \frac{\partial D}{\partial G} \frac{\partial G}{\partial \theta_G} \\ &= -\frac{1}{1} \sigma'(v) \frac{\partial v}{\partial G} \frac{\partial G}{\partial \theta_G} \\ &\approx -1 \times 0 \times \frac{\partial v}{\partial G} \frac{\partial G}{\partial \theta_G}\end{aligned}$$


Closer look at f_{acc}

- To accelerate training early on, we can instead use a different loss term for the fake data that yields the same desired behavior but trains faster.

$$f_{\text{acc}}(\theta_G, \theta_D) = \mathbb{E}_{\mathbf{x} \sim P_{\text{data}}(\mathbf{x})} [\log D_{\theta_D}(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim P(\mathbf{z})} [-\log(D_{\theta_D}(G_{\theta_G}(\mathbf{z})))]$$

New loss term

- The reason is that the gradient of $-\log(v)$ for $v \approx 0$ is very large, whereas the gradient of $\log(1-v) \approx 1$.




GANs

- Show Goodfellow et al. 2014 paper.
- GANs represent the state-of-the-art (as of 2021) for generating realistic data.
- GANs have also inspired many other adversarial training methods.

Conditional GANs

- One example is **conditional GANs**:
 - G also accepts a parameter vector \mathbf{f} (e.g., 1-hot encoding of MNIST class) that specifies what *kind* of data to generate.
 - D also accepts \mathbf{f} to help discriminate a particular kind of real from fake data.



Difficulty in training

- GANs are renowned for being difficult to train:
 1. How to choose k , l ? More hyperparameters to optimize.

Difficulty in training

- GANs are renowned for being difficult to train:
 1. How to choose k , l ? More hyperparameters to optimize.
 2. We will probably never reach the equilibrium where G exactly produces $P_{\text{data}}(\mathbf{x})$ and D 's accuracy is 0.5.
 - What kind of “training curve” for D , G should we expect?
 - If D gets too good too fast, then G may never have a chance to improve.

Difficulty in training

- GANs are renowned for being difficult to train:

3. Mode collapse — G generates realistic data but only for a *subset* of the domain of $P_{\text{data}}(\mathbf{x})$, e.g.:





Difficulty in training

- GANs are renowned for being difficult to train:

3. Neuron co-adaptation — training gets stuck because multiple pathways rely on each other too much.

- Consider an MNIST image near the borders: what property do pixels x_1, x_2 have?

- $x_1 = x_2 = 0$?





Difficulty in training

- GANs are renowned for being difficult to train:

3. Neuron co-adaptation — training gets stuck because multiple pathways rely on each other too much.

- Consider an MNIST image near the borders: what property do pixels x_1, x_2 have?

- $ax_1 - ax_2 = 0$?





Difficulty in training

- GANs are renowned for being difficult to train:

3. Neuron co-adaptation — training gets stuck because multiple pathways rely on each other too much.

- In the latter case, D gives feedback to G that images are “ok” as long as the background noise “cancels” itself, e.g.:





Difficulty in training

- GANs are renowned for being difficult to train:

3. Neuron co-adaptation — training gets stuck because multiple pathways rely on each other too much.

- In the latter case, D gives feedback to G that images are “ok” as long as the background noise “cancels” itself, e.g.:





Difficulty in training

- GANs are renowned for being difficult to train:

3. Neuron co-adaptation — training gets stuck because multiple pathways rely on each other too much.

- In the latter case, D gives feedback to G that images are “ok” as long as the background noise “cancels” itself, e.g.:





Difficulty in training

- GANs are renowned for being difficult to train:

3. Neuron co-adaptation — training gets stuck because multiple pathways rely on each other too much.

- To prevent this from occurring, we can use dropout on the input layer \mathbf{x} , so that each pathway must judge independently if the image is a fake.



Adversarial training examples

Adversarial training examples

- One of the weaknesses of DL is that we often do not know how the models work.
- This means that models can sometimes be systematically exploited to give nonsensical outputs.
- In high-stakes scenarios such as autonomous vehicles, the results can be disastrous (e.g., misread a speed limit sign of “25km/h” for “95km/h”).

Adversarial training examples

- One way to systematically confuse a trained NN is to “adjust” an input \mathbf{x} so as to *maximally change* the output \hat{y} .
- In other words, for any loss function f (e.g., cross-entropy), compute the gradient:

$$\nabla_{\mathbf{x}} f(\mathbf{x}; \mathbf{w})$$

where \mathbf{w} are the NN's weights.

- We then construct an **adversarial example** (that differs by at most η from \mathbf{x}) as:

$$\mathbf{x}' = \mathbf{x} + \eta \text{sign}(\nabla_{\mathbf{x}} f(\mathbf{x}; \mathbf{w}))$$

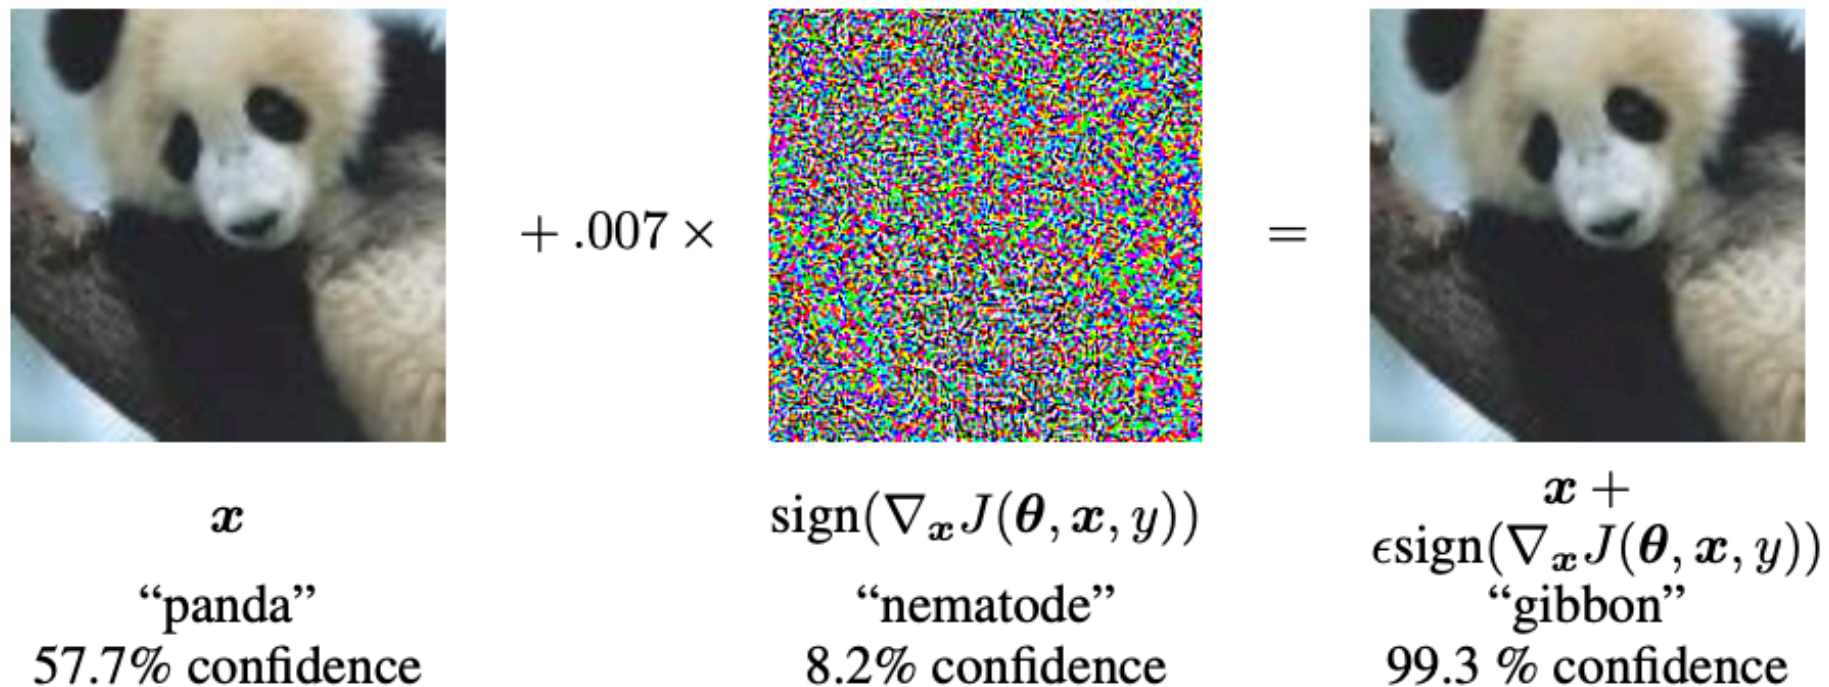
Adversarial training examples

- Note that this is *different* from what we usually do during training, i.e., conduct SGD on \mathbf{w} using the gradient:

$$\nabla_{\mathbf{w}} f(\mathbf{x}; \mathbf{w})$$

Adversarial training examples

- This technique is responsible for the famous panda -> gibbon mistake (Goodfellow et al. 2014):



x
“panda”
57.7% confidence

$+ .007 \times$

$\text{sign}(\nabla_x J(\theta, x, y))$
“nematode”
8.2% confidence

$=$

$x + \epsilon \text{sign}(\nabla_x J(\theta, x, y))$
“gibbon”
99.3 % confidence

Adversarial training examples

- Since this discovery of this attack, many DL researchers have investigated how to prevent it.
- One simple technique that can help is to *train* explicitly on adversarial examples.
- This is easy, since we can create adversarial examples at will, and can improve robustness to outliers at test time.